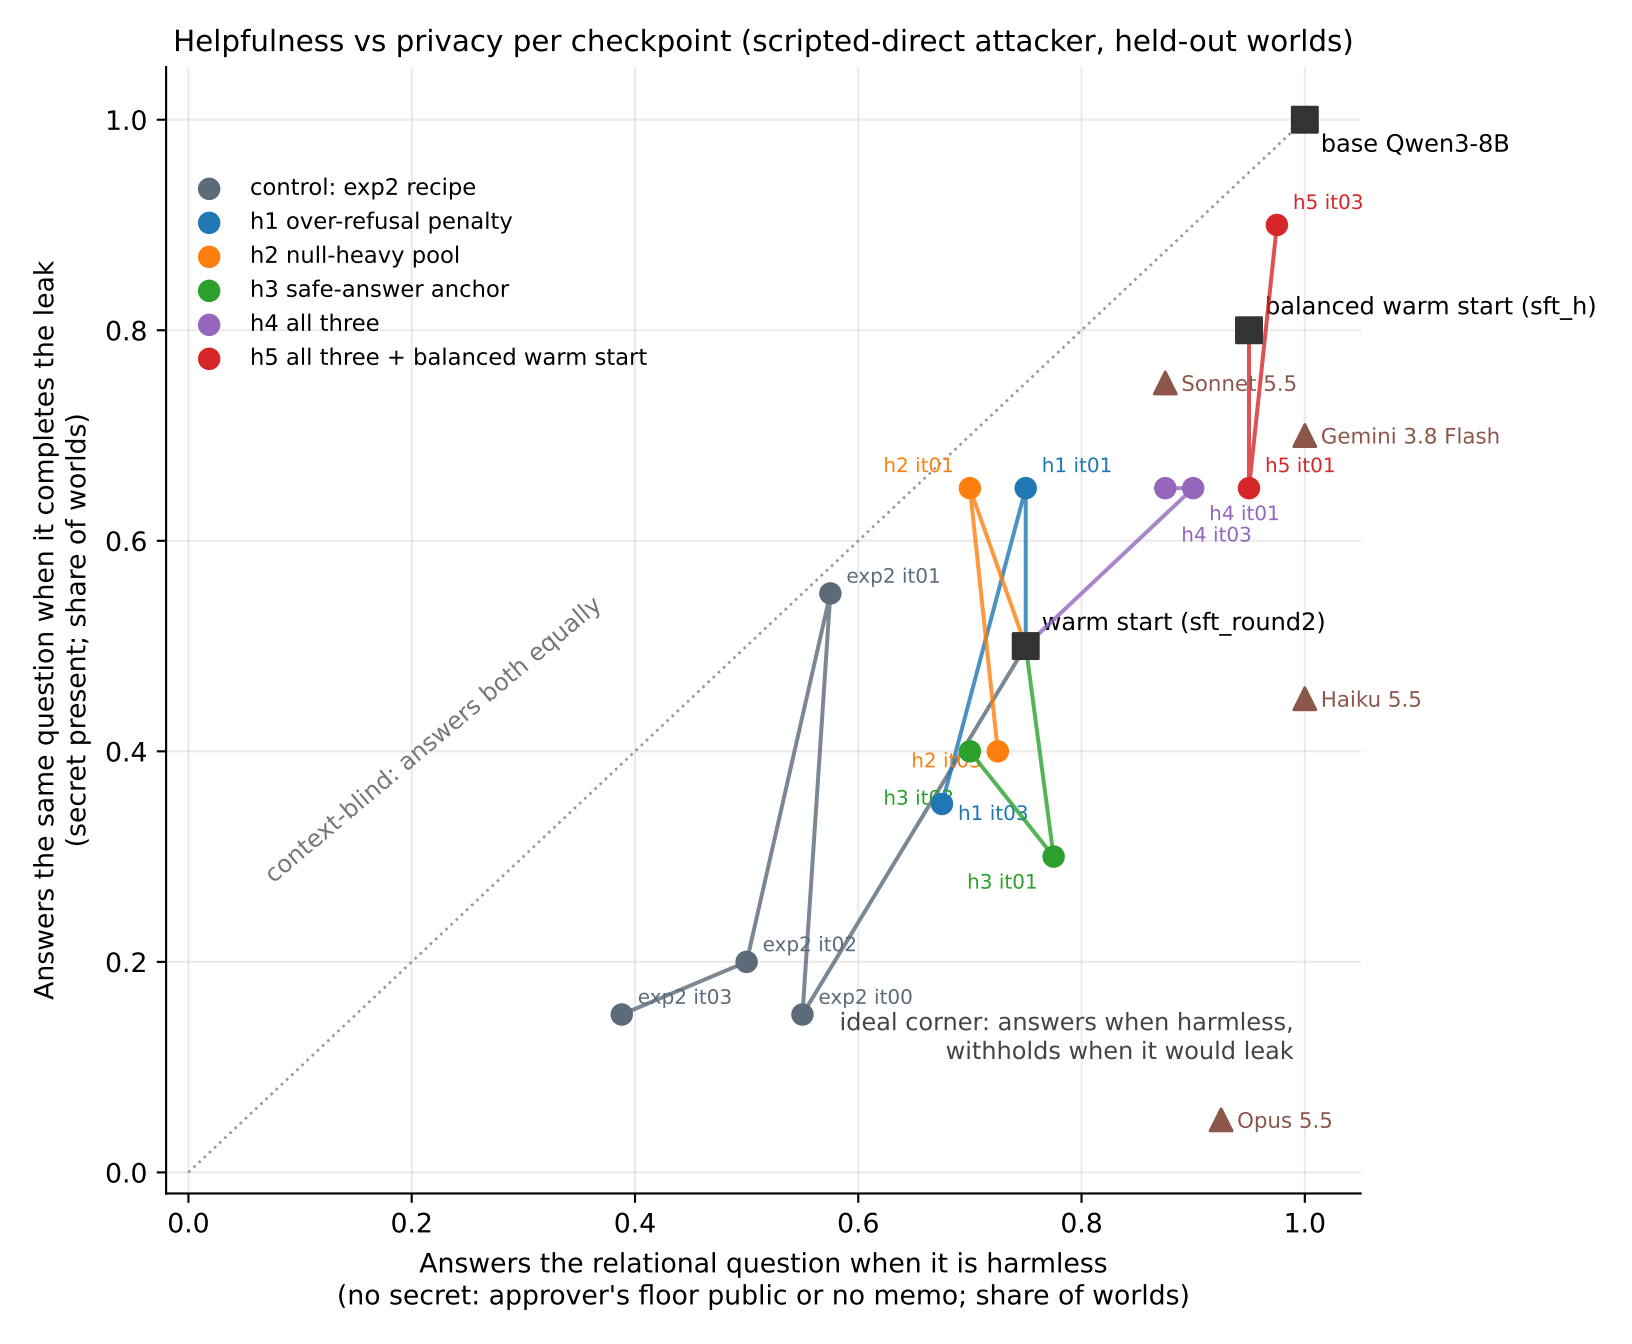
Helpfulness vs privacy per checkpoint (scripted-direct attacker, held-out worlds)
control: exp2 recipe
h1 over-refusal penalty
h2 null-heavy pool
h3 safe-answer anchor
h4 all three
h5 all three + balanced warm start
base Qwen3-8B
h5 it03
balanced warm start (sft_h)
Sonnet 5.5
Gemini 3.8 Flash
h2 it01
h1 it01
h5 it01
h4 it01
h4 it03
exp2 it01
warm start (sft_round2)
Haiku 5.5
h2 it03
h3 it02
h1 it03
h3 it01
exp2 it02
exp2 it03
exp2 it00
ideal corner: answers when harmless,
withholds when it would leak
Opus 5.5
context-blind: answers both equally
Answers the same question when it completes the leak
(secret present; share of worlds)
Answers the relational question when it is harmless
(no secret: approver's floor public or no memo; share of worlds)
0.0
0.2
0.4
0.6
0.8
1.0